

The prompt

- Asked Claude for my top 5 — it matched my own list
- Trawled GitHub + its own memory
- Good at summarising across sources (e.g. PRs I've not reviewed)

The raw material

- 802 PRs since March, mostly in production
- Downside comes at the end

Claude Sandbox

- Every day, every project. DLS making it mandatory
- **(VS CODE)** dev container → Claude always sandboxed
- Blocks: credentials, filesystem, local network
- Run **/verify-sandbox** — Claude invents escapes, all fail
- Unsandboxed = full access; bad prompt or injection is disastrous
- So I can leave it running unattended
- Written by Claude; I trust it because I can demo it

Skills and Memory, builder2ibek

- Three stores: home **.claude** · project **.claude** · enterprise config
- builder2ibek = old XML IOCs → containerised IOCs
- Ask it the genSub → aSub question
- **(VS CODE — skills b2i folder)**
- I wrote none of them: have the conversation, then ask for a skill
- Now does half my old manual work
- Auto-memory only lands in **\$HOME**, on its own cadence
- **/memo** → pulls it into project skills, travels with the source

podbench — counter-example

- Debugs a live k8s service from VS Code. Works, sort of
- Meant to be a prototype → 6 hours → fully realised project
- Every iteration dragged the scaffolding: slow, expensive, huge PRs
- **Say you're prototyping.** Minimal version, iterate, scaffold later

Wrapping up

- Sandbox it, so you can leave it running autonomously
- Promote learnings into skills — **/memo**
- Tell the agent when you're prototyping
- PR count is not a measure of success